

## Part-2

Florencia Luque and Seyed Amirhossein Mosaddad

2024-10-18

### Part 2

```
library(caret)
library(caretEnsemble)
library(h2o)
library(tidymodels)
library(ROCR)
library(ConfusionTableR)
library(dplyr)
library(pander)

path = "/Users/soroush/Desktop/UC3M/courses/R/minitask1/Stroke_data"
stroke_data = read.csv("healthcare-dataset-stroke-data.csv", header = TRUE)

stroke_data = na.omit(stroke_data)
stroke_data = stroke_data[stroke_data$gender!="Other",]
stroke_data$gender = as.factor(stroke_data$gender)
stroke_data$hypertension = as.factor(stroke_data$hypertension)
stroke_data$heart_disease = as.factor(stroke_data$heart_disease)
stroke_data$work_type = as.factor(stroke_data$work_type)
stroke_data$Residence_type = as.factor(stroke_data$Residence_type)
stroke_data$ever_married = as.factor(stroke_data$ever_married)
stroke_data$smoking_status = as.factor(stroke_data$smoking_status)
stroke_data$stroke = as.factor(stroke_data$stroke)
stroke_data$bmi = as.numeric(stroke_data$bmi)

frec_table <- cut(stroke_data$bmi, breaks = c(0, 18.5, 24.9, 29.9, 34.9, Inf),
                  labels = c("Underweight", "Normal", "Overweight", "Obese", "Extremely Obese"))
stroke_data$cat_weight = frec_table

glucose_categories <- cut(stroke_data$avg_glucose_level,
                          breaks = c(0, 99.9, 125.9, Inf),
                          labels = c("Normal", "Prediabetes", "Diabetes"))

# Reorder the levels to put Normal in the middle
glucose_categories <- factor(glucose_categories, levels = c("Prediabetes", "Normal", "Diabetes"))

# Add the categories to the data frame
stroke_data$glucose_category = glucose_categories
str(stroke_data)
```

```
## 'data.frame':    5109 obs. of  14 variables:
## $ id            : int  9046 51676 31112 60182 1665 56669 53882 10434 27419 60491 ...
## $ gender        : Factor w/ 2 levels "Female","Male": 2 1 2 1 1 2 2 1 1 1 ...
## $ age           : num  67 61 80 49 79 81 74 69 59 78 ...
## $ hypertension  : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 2 1 1 1 ...
## $ heart_disease : Factor w/ 2 levels "0","1": 2 1 2 1 1 1 2 1 1 1 ...
## $ ever_married  : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 2 2 1 2 2 ...
## $ work_type     : Factor w/ 5 levels "children","Govt_job",...: 4 5 4 4 5 4 4 4 4 4 ...
## $ Residence_type : Factor w/ 2 levels "Rural","Urban": 2 1 1 2 1 2 1 2 1 2 ...
## $ avg_glucose_level: num  229 202 106 171 174 ...
## $ bmi           : num  36.6 NA 32.5 34.4 24 29 27.4 22.8 NA 24.2 ...
## $ smoking_status : Factor w/ 4 levels "formerly smoked",...: 1 2 2 3 2 1 2 2 4 4 ...
## $ stroke        : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
## $ cat_weight     : Factor w/ 5 levels "Underweight",...: 5 NA 4 4 2 3 3 2 NA 2 ...
## $ glucose_category : Factor w/ 3 levels "Prediabetes",...: 3 3 1 3 3 3 2 2 2 2 ...
```

## Relation between variables

```
tab_st_gd = table(stroke_data$stroke,stroke_data$gender)
chisq.test(tab_st_gd)
```

### Stroke with Gender

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_st_gd
## X-squared = 0.34, df = 1, p-value = 0.5598
```

The *p-value* is larger than 0.05 this mean that there's not evidence of dependency between the variables gender and stroke.

```
tab_st_hy = table(stroke_data$stroke,stroke_data$hypertension)
chisq.test(tab_st_hy)
```

### Stroke with Hypertension

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_st_hy
## X-squared = 81.573, df = 1, p-value < 2.2e-16
```

The *p-value* is a lot smaller than 0.05. This mean that there's a relation between getting a stroke and hypertension. This is a could be a comprobaton of the hypothesis that we established earlier about the existence of a relation between this two variables.

```
tab_st_hd = table(stroke_data$stroke,stroke_data$heart_disease)
chisq.test(tab_st_hd)
```

### Stroke with Heart Disease

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_st_hd
## X-squared = 90.229, df = 1, p-value < 2.2e-16
```

The *p-value* is a lot smaller than 0.05. This mean that there's a relation between getting a stroke and heart disease This is a could be a comprobaton of the hypothesis that we established earlier about the existence of a relation between this two variables.

```
tab_st_rt = table(stroke_data$stroke,stroke_data$Residence_type)
chisq.test(tab_st_rt)
```

### Stroke with Residence Type

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_st_rt
## X-squared = 1.075, df = 1, p-value = 0.2998
```

As we had seen in the graph there's no evidence to say that there's a relation between the type of residence and getting a stroke.

```
tab_st_em = table(stroke_data$stroke,stroke_data$ever_married)
chisq.test(tab_st_em)
```

### Stroke with Ever Married

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_st_em
## X-squared = 58.868, df = 1, p-value = 1.686e-14
```

Apparently there's a relation within this two variables. Having a stroke have a relation with have been or had been ever married.

```
tab_st_sk = table(stroke_data$stroke,stroke_data$smoking_status)
chisq.test(tab_st_sk)
```

### Stroke with Smoking Status

```
##
## Pearson's Chi-squared test
##
## data:  tab_st_sk
## X-squared = 29.226, df = 3, p-value = 2.008e-06
```

The *p-value* is a lot smaller than 0.05 so there's a relation between the variables. This was something that we *don't know what to write here*

```
tab_st_wt = table(stroke_data$stroke,stroke_data$work_type)
chisq.test(tab_st_wt)
```

### Stroke with Work Type

```
##
## Pearson's Chi-squared test
##
## data:  tab_st_wt
## X-squared = 49.159, df = 4, p-value = 5.409e-10
```

There's a relation between the variables ( $p\text{-value} < 0.05$ ). This we think was because of the difference between the quantity of people who got a stroke and work independently and the people who work with children because the difference was big between them.

```
tab_hy_hd = table(stroke_data$heart_disease,stroke_data$hypertension)
chisq.test(tab_hy_hd)
```

### Heart Disease and Hypertension

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_hy_hd
## X-squared = 58.31, df = 1, p-value = 2.239e-14
```

There's a relation between hypertension and heart disease and both variable are related to stroke. This could be a good indicator that within only one of this variables we could have the same information in the model.

```
tab_hd_em = table(stroke_data$heart_disease,stroke_data$ever_married)
chisq.test(tab_hd_em)
```

### Heart Disease and Ever Married

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_hd_em
## X-squared = 66.036, df = 1, p-value = 4.428e-16
```

```
tab_hd_sk= table(stroke_data$heart_disease,stroke_data$smoking_status)
chisq.test(tab_hd_sk)
```

### Heart Disease and Smoking Status

```
##
## Pearson's Chi-squared test
##
## data:  tab_hd_sk
## X-squared = 44.74, df = 3, p-value = 1.051e-09
```

```
tab_hd_wt = table(stroke_data$heart_disease,stroke_data$work_type)
chisq.test(tab_hd_wt)
```

### Heart Disease and Work Type

```
##
## Pearson's Chi-squared test
##
## data:  tab_hd_wt
## X-squared = 70.689, df = 4, p-value = 1.623e-14
```

```
tab_hy_em = table(stroke_data$heart_disease,stroke_data$ever_married)
chisq.test(tab_hy_em)
```

### Hypertension and Ever Married

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  tab_hy_em
## X-squared = 66.036, df = 1, p-value = 4.428e-16
```

```
tab_hy_sk = table(stroke_data$heart_disease,stroke_data$smoking_status)
chisq.test(tab_hy_sk)
```

## Hypertension and Smoking Status

```
##
## Pearson's Chi-squared test
##
## data:  tab_hy_sk
## X-squared = 44.74, df = 3, p-value = 1.051e-09
```

```
tab_hy_wt= table(stroke_data$heart_disease,stroke_data$work_type)
chisq.test(tab_hy_wt)
```

## Hypertension and Work Type

```
##
## Pearson's Chi-squared test
##
## data:  tab_hy_wt
## X-squared = 70.689, df = 4, p-value = 1.623e-14
```

## Stroke with Age

```
t.test(age ~ stroke, data = stroke_data)
```

```
##
## Welch Two Sample t-test
##
## data:  age by stroke
## t = -29.682, df = 331.68, p-value < 2.2e-16
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
##  -27.46015 -24.04658
## sample estimates:
## mean in group 0 mean in group 1
##      41.97483      67.72819
```

As you can see we reject the null hypothesis that said that both group have the same mean and this tell as that this variable could have and impact in the probability of getting a stroke in this case getting older increase your chances.

## Stroke with Average Glucose Levels

```
t.test(avg_glucose_level ~ stroke, data = stroke_data)
```

```
##
## Welch Two Sample t-test
##
## data: avg_glucose_level by stroke
## t = -6.9844, df = 260.9, p-value = 2.373e-11
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## -35.58269 -19.93162
## sample estimates:
## mean in group 0 mean in group 1
## 104.7876 132.5447
```

Also we can say that the difference in means between the groups with a stroke are without is not zero.

## Stroke with BMI

```
t.test(bmi ~ stroke, data = stroke_data)
```

```
##
## Welch Two Sample t-test
##
## data: bmi by stroke
## t = -3.6374, df = 237.84, p-value = 0.0003377
## alternative hypothesis: true difference in means between group 0 and group 1 is not equal to 0
## 95 percent confidence interval:
## -2.5387991 -0.7549231
## sample estimates:
## mean in group 0 mean in group 1
## 28.82443 30.47129
```

Apparently all the continuous variables could have an impact in the chances of getting a stroke. This variable *bmi* also has a *p-value* less than 0.05 so we can say that the groups have significant different means.

## Models

### Logistic Regression (Flo)

```
set.seed(23)
index_split = createDataPartition(stroke_data$stroke,p=0.8,list=FALSE)
train = stroke_data[index_split,]
test = stroke_data[-index_split,]
train = subset(train,select = -c(id,cat_weight,glucose_category))
train = na.omit(train)
test = subset(test,select = -c(id,cat_weight,glucose_category))
test = na.omit(test)
```

We would start the models with one that includes all the variables and then reduce it with different tests.

```
class_weight <- ifelse(train$stroke == 1, 25, 1.04)
log_reg_model = glm(stroke~(gender + age + hypertension + heart_disease + ever_married +
  work_type + Residence_type + avg_glucose_level + bmi + smoking_status),data = train,family = binomial,
summary(log_reg_model)
```

```
##
## Call:
## glm(formula = stroke ~ (gender + age + hypertension + heart_disease +
##   ever_married + work_type + Residence_type + avg_glucose_level +
##   bmi + smoking_status), family = binomial(link = "logit"),
##   data = train, weights = class_weight)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.860e+00  2.485e-01 -15.532  < 2e-16 ***
## genderMale      -5.040e-02  6.005e-02  -0.839  0.401304
## age             7.914e-02  2.321e-03  34.104  < 2e-16 ***
## hypertension1    8.490e-01  7.759e-02  10.942  < 2e-16 ***
## heart_disease1   3.797e-01  1.003e-01   3.784  0.000154 ***
## ever_marriedYes -1.144e-01  9.289e-02  -1.231  0.218197
## work_typeGovt_job -1.359e+00  2.611e-01  -5.204  1.95e-07 ***
## work_typeNever_worked -1.123e+01  1.329e+02  -0.084  0.932663
## work_typePrivate -1.292e+00  2.549e-01  -5.070  3.98e-07 ***
## work_typeSelf-employed -1.552e+00  2.679e-01  -5.793  6.91e-09 ***
## Residence_typeUrban  3.577e-02  5.762e-02   0.621  0.534725
## avg_glucose_level  3.118e-03  5.853e-04   5.326  1.00e-07 ***
## bmi             1.599e-02  4.514e-03   3.543  0.000395 ***
## smoking_statusnever smoked -3.382e-01  7.671e-02  -4.409  1.04e-05 ***
## smoking_statussmokes  2.248e-01  8.857e-02   2.538  0.011138 *
## smoking_statusUnknown -5.600e-01  9.272e-02  -6.040  1.54e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 11203.1  on 3928  degrees of freedom
## Residual deviance:  7500.3  on 3913  degrees of freedom
## AIC: 7377.4
##
## Number of Fisher Scoring iterations: 12
```

```
fitted.results <- predict(log_reg_model,newdata = test,type='response')
fitted.results <- ifelse(fitted.results > 0.8,1,0)
test$accu = fitted.results
misClasificError <- mean(fitted.results != test$stroke)
print(paste('Accuracy',1-misClasificError))
```

```
## [1] "Accuracy 0.892747701736466"
```



```
#pander(confusionMatrix(test$accu, test$stroke))
```

As you can see if we take 0.8 as the threshold we get an accuracy of 0.89 and a sensibility of 0.9. We will continue deleting the variable resident type because it's not have any significance in the model.

```
class_weight <- ifelse(train$stroke == 1, 25, 1.04)
log_reg_model2 = glm(stroke~(gender + age + hypertension + heart_disease + ever_married +work_type + avg
summary(log_reg_model2)
```

```
##
## Call:
## glm(formula = stroke ~ (gender + age + hypertension + heart_disease +
##     ever_married + work_type + avg_glucose_level + bmi + smoking_status),
##     family = binomial(link = "logit"), data = train, weights = class_weight)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.845e+00  2.474e-01 -15.543  < 2e-16 ***
## genderMale      -5.104e-02  6.004e-02  -0.850  0.395203
## age             7.919e-02  2.319e-03  34.147  < 2e-16 ***
## hypertension1    8.483e-01  7.758e-02  10.935  < 2e-16 ***
## heart_disease1    3.795e-01  1.004e-01   3.781  0.000156 ***
## ever_marriedYes  -1.188e-01  9.254e-02  -1.284  0.199235
## work_typeGovt_job -1.355e+00  2.610e-01  -5.192  2.08e-07 ***
## work_typeNever_worked -1.122e+01  1.329e+02  -0.084  0.932721
## work_typePrivate  -1.288e+00  2.548e-01  -5.057  4.26e-07 ***
## work_typeSelf-employed -1.549e+00  2.679e-01  -5.784  7.29e-09 ***
## avg_glucose_level  3.103e-03  5.849e-04   5.306  1.12e-07 ***
## bmi             1.609e-02  4.509e-03   3.569  0.000359 ***
## smoking_statusnever smoked -3.382e-01  7.673e-02  -4.408  1.05e-05 ***
## smoking_statussmokes  2.259e-01  8.853e-02   2.552  0.010720 *
## smoking_statusUnknown -5.592e-01  9.272e-02  -6.031  1.63e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 11203.1  on 3928  degrees of freedom
## Residual deviance:  7500.7  on 3914  degrees of freedom
## AIC: 7375.8
##
## Number of Fisher Scoring iterations: 12
```

```
fitted.results <- predict(log_reg_model2,newdata = test,type='response')
fitted.results <- ifelse(fitted.results > 0.8,1,0)
test$accu = fitted.results
misClasificError <- mean(fitted.results != test$stroke)
print(paste('Accuracy',1-misClasificError))
```

```
## [1] "Accuracy 0.892747701736466"
```

```
pander(confusionMatrix(as.factor(test$accu),test$stroke))
```

- **positive:** 0
- **table:**

	0	1
<b>0</b>	851	19
<b>1</b>	86	23

- **overall:**

Table 2: Table continues below

Accuracy	Kappa	AccuracyLower	AccuracyUpper	AccuracyNull
0.8927	0.2587	0.8717	0.9114	0.9571

AccuracyPValue	McnemarPValue
1	1.187e-10

- **byClass:**

Table 4: Table continues below

Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision
0.9082	0.5476	0.9782	0.211	0.9782

Table 5: Table continues below

Recall	F1	Prevalence	Detection Rate	Detection Prevalence
0.9082	0.9419	0.9571	0.8693	0.8887

Balanced Accuracy
0.7279

- **mode:** sens\_spec
- **dots:**

The AIC was reduce so the model improved just a little.

```
anova(log_reg_model2,log_reg_model)
```

```
## Analysis of Deviance Table
##
## Model 1: stroke ~ (gender + age + hypertension + heart_disease + ever_married +
##   work_type + avg_glucose_level + bmi + smoking_status)
## Model 2: stroke ~ (gender + age + hypertension + heart_disease + ever_married +
##   work_type + Residence_type + avg_glucose_level + bmi + smoking_status)
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1      3914      7500.7
## 2      3913      7500.3  1  0.38543   0.5347
```

As the *p-value* is higher than 0.05 we cannot reject that the simpler model is better. So we are going to deleted the variable gender because i doesn't affect the model.

```
class_weight <- ifelse(train$stroke == 1, 25, 1.04)
log_reg_model3 = glm(stroke~(ever_married + age + hypertension + heart_disease+work_type + avg_glucose_level),
summary(log_reg_model3)
```

```
##
## Call:
## glm(formula = stroke ~ (ever_married + age + hypertension + heart_disease +
##   work_type + avg_glucose_level + bmi + smoking_status), family = binomial(link = "logit"),
##   data = train, weights = class_weight)
##
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -3.871e+00  2.456e-01 -15.761  < 2e-16 ***
## ever_marriedYes    -1.170e-01  9.243e-02  -1.266  0.205554
## age                7.913e-02  2.318e-03  34.146  < 2e-16 ***
## hypertension1      8.493e-01  7.758e-02  10.947  < 2e-16 ***
## heart_disease1     3.682e-01  9.943e-02   3.703  0.000213 ***
## work_typeGovt_job  -1.347e+00  2.608e-01  -5.163  2.43e-07 ***
## work_typeNever_worked -1.122e+01  1.330e+02  -0.084  0.932782
## work_typePrivate    -1.283e+00  2.547e-01  -5.038  4.71e-07 ***
## work_typeSelf-employed -1.542e+00  2.677e-01  -5.759  8.48e-09 ***
## avg_glucose_level   3.073e-03  5.837e-04   5.264  1.41e-07 ***
## bmi                1.612e-02  4.508e-03   3.576  0.000349 ***
## smoking_statusnever smoked -3.298e-01  7.609e-02  -4.334  1.46e-05 ***
## smoking_statussmokes  2.259e-01  8.857e-02   2.550  0.010762 *
## smoking_statusUnknown  -5.577e-01  9.267e-02  -6.019  1.76e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##   Null deviance: 11203.1  on 3928  degrees of freedom
## Residual deviance: 7501.4  on 3915  degrees of freedom
## AIC: 7374.5
##
## Number of Fisher Scoring iterations: 12

fitted.results <- predict(log_reg_model3,newdata = test,type='response')
fitted.results <- ifelse(fitted.results > 0.8,1,0)
test$accu = fitted.results
```

```
misClassificError <- mean(fitted.results != test$stroke)
print(paste('Accuracy', 1-misClassificError))
```

```
## [1] "Accuracy 0.889683350357508"
```

```
pander(confusionMatrix(as.factor(test$accu), test$stroke))
```

- **positive:** 0
- **table:**

	0	1
0	848	19
1	89	23

- **overall:**

Table 8: Table continues below

Accuracy	Kappa	AccuracyLower	AccuracyUpper	AccuracyNull
0.8897	0.252	0.8684	0.9086	0.9571

AccuracyPValue	McnemarPValue
1	3.147e-11

- **byClass:**

Table 10: Table continues below

Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision
0.905	0.5476	0.9781	0.2054	0.9781

Table 11: Table continues below

Recall	F1	Prevalence	Detection Rate	Detection Prevalence
0.905	0.9401	0.9571	0.8662	0.8856

Balanced Accuracy
0.7263

- **mode:** sens\_spec
- **dots:**

```
anova(log_reg_model3,log_reg_model2)
```

```
## Analysis of Deviance Table
##
## Model 1: stroke ~ (ever_married + age + hypertension + heart_disease +
##   work_type + avg_glucose_level + bmi + smoking_status)
## Model 2: stroke ~ (gender + age + hypertension + heart_disease + ever_married +
##   work_type + avg_glucose_level + bmi + smoking_status)
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1      3915      7501.4
## 2      3914      7500.7  1  0.72316   0.3951
```

As the *p-value* is higher than 0.05 we cannot reject that the simpler model is better. So we are going to deleted the variable ever married because i doesn't affect the model.

```
class_weight <- ifelse(train$stroke == 1, 25, 1.04)
log_reg_model4 = glm(stroke~(age + hypertension + heart_disease+work_type + avg_glucose_level + bmi + s
summary(log_reg_model4)
```

```
##
## Call:
## glm(formula = stroke ~ (age + hypertension + heart_disease +
##   work_type + avg_glucose_level + bmi + smoking_status), family = binomial(link = "logit"),
##   data = train, weights = class_weight)
##
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.861e+00  2.455e-01 -15.729  < 2e-16 ***
## age              7.861e-02  2.284e-03  34.412  < 2e-16 ***
## hypertension1    8.531e-01  7.740e-02  11.023  < 2e-16 ***
## heart_disease1   3.719e-01  9.928e-02   3.746 0.000180 ***
## work_typeGovt_job -1.417e+00  2.556e-01  -5.544 2.96e-08 ***
## work_typeNever_worked -1.120e+01  1.323e+02  -0.085 0.932489
## work_typePrivate  -1.355e+00  2.490e-01  -5.443 5.24e-08 ***
## work_typeSelf-employed -1.615e+00  2.621e-01  -6.165 7.07e-10 ***
## avg_glucose_level  3.020e-03  5.822e-04   5.186 2.14e-07 ***
## bmi              1.580e-02  4.505e-03   3.507 0.000454 ***
## smoking_statusnever smoked -3.190e-01  7.559e-02  -4.220 2.44e-05 ***
## smoking_statussmokes  2.288e-01  8.858e-02   2.583 0.009806 **
## smoking_statusUnknown  -5.505e-01  9.252e-02  -5.950 2.68e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##   Null deviance: 11203  on 3928  degrees of freedom
## Residual deviance:  7503  on 3916  degrees of freedom
## AIC: 7374.1
##
## Number of Fisher Scoring iterations: 12
```

```
fitted.results <- predict(log_reg_model4,newdata = test,type='response')
fitted.results <- ifelse(fitted.results > 0.8,1,0)
test$accu = fitted.results
misClasificError <- mean(fitted.results != test$stroke)
print(paste('Accuracy',1-misClasificError))
```

```
## [1] "Accuracy 0.888661899897855"
```

```
pander(confusionMatrix(as.factor(test$accu),test$stroke))
```

- **positive:** 0
- **table:**

	0	1
0	848	20
1	89	22

- **overall:**

Table 14: Table continues below

Accuracy	Kappa	AccuracyLower	AccuracyUpper	AccuracyNull
0.8887	0.2403	0.8673	0.9077	0.9571

AccuracyPValue	McnemarPValue
1	7.356e-11

- **byClass:**

Table 16: Table continues below

Sensitivity	Specificity	Pos Pred Value	Neg Pred Value	Precision
0.905	0.5238	0.977	0.1982	0.977

Table 17: Table continues below

Recall	F1	Prevalence	Detection Rate	Detection Prevalence
0.905	0.9396	0.9571	0.8662	0.8866

Balanced Accuracy
0.7144

- **mode:** sens\_spec
- **dots:**

```
anova(log_reg_model,log_reg_model4)
```

```
## Analysis of Deviance Table
##
## Model 1: stroke ~ (gender + age + hypertension + heart_disease + ever_married +
##   work_type + Residence_type + avg_glucose_level + bmi + smoking_status)
## Model 2: stroke ~ (age + hypertension + heart_disease + work_type + avg_glucose_level +
##   bmi + smoking_status)
##   Resid. Df Resid. Dev Df Deviance Pr(>Chi)
## 1      3913      7500.3
## 2      3916      7503.0 -3   -2.7142   0.4378
```

The simpler model is better so as all the variables are significance in the model we will leave it at that.

```
pander(exp(log_reg_model4$coefficients))
```

Table 19: Table continues below

(Intercept)	age	hypertension1	heart_disease1	work_typeGovt_job
0.02104	1.082	2.347	1.45	0.2425

Table 20: Table continues below

work_typeNever_worked	work_typePrivate	work_typeSelf-employed
1.362e-05	0.2579	0.1988

Table 21: Table continues below

avg_glucose_level	bmi	smoking_statusnever smoked	smoking_statussmokes
1.003	1.016	0.7269	1.257

  

smoking_statusUnknown
0.5767

With the exp of the coefficients we get the odds of getting a stroke. Like the intercept means that if all the other variables are zero you have a 2.1% odds of getting a stroke. As we said in the preliminary analysis if you got hypertension the odd of getting a stroke duplicate. If you got a heart disease you have a 4.5% increase in your odds.

**VSM (soroush)**

**XGBoost for classification (soroush)**

**H2O models with autoh2o (Flo)**

Split the data to get a 80% for training and 20% for testing. #flo

```
h2o.init()
```

```
## Connection successful!
##
## R is connected to the H2O cluster:
##   H2O cluster uptime:      1 days 3 hours
##   H2O cluster timezone:    Europe/Paris
##   H2O data parsing timezone: UTC
##   H2O cluster version:     3.44.0.3
##   H2O cluster version age:  10 months and 4 days
##   H2O cluster name:        H2O_started_from_R_flore_dfs870
##   H2O cluster total nodes: 1
##   H2O cluster total memory: 3.49 GB
##   H2O cluster total cores: 16
##   H2O cluster allowed cores: 16
##   H2O cluster healthy:     TRUE
##   H2O Connection ip:       localhost
##   H2O Connection port:     54321
##   H2O Connection proxy:    NA
##   H2O Internal Security:   FALSE
##   R Version:                R version 4.4.1 (2024-06-14 ucrt)
```

```
stroke_h2o=as.h2o(stroke_data)
```

```
## |
```

```
split_data = h2o.splitFrame(data=stroke_h2o,ratios=0.8,seed=23)
train_h2o = split_data[[1]]
test_h2o = split_data[[2]]
predictor = c("gender","age","hypertension","heart_disease","ever_married","work_type","Residence_type")
aml = h2o.automl(x = predictor,y="stroke",training_frame=train_h2o,max_models=10,keep_cross_validation_
```

```
## |
## 13:04:31.75: AutoML: XGBoost is not available; skipping it. |
```

```
lb <- aml@leaderboard
print(lb, n = nrow(lb),extra_columns="ALL")
```

```
##                               model_id      auc  logloss
## 1   StackedEnsemble_AllModels_1_AutoML_13_20241025_130431 0.8325502 0.1609698
## 2   StackedEnsemble_BestOfFamily_1_AutoML_13_20241025_130431 0.8314759 0.1611657
## 3                               GLM_1_AutoML_13_20241025_130431 0.8295080 0.1619212
## 4                               GBM_1_AutoML_13_20241025_130431 0.8268580 0.1678332
## 5                               XRT_1_AutoML_13_20241025_130431 0.8232607 0.1928370
## 6                               GBM_2_AutoML_13_20241025_130431 0.8148693 0.1841245
## 7                               GBM_5_AutoML_13_20241025_130431 0.8143710 0.1762200
## 8                               GBM_3_AutoML_13_20241025_130431 0.8073816 0.1885697
## 9                               GBM_4_AutoML_13_20241025_130431 0.8028584 0.2025739
## 10                              DRF_1_AutoML_13_20241025_130431 0.7935992 0.3196497
## 11   GBM_grid_1_AutoML_13_20241025_130431_model_1 0.7633213 0.2102651
## 12   DeepLearning_1_AutoML_13_20241025_130431 0.7041979 6.9384679
```



```
##          aucpr mean_per_class_error      rmse      mse
## 1  0.16823965          0.2918846 0.2083456 0.04340791
## 2  0.16656006          0.2648105 0.2083859 0.04342470
## 3  0.17453868          0.2979178 0.2084133 0.04343609
## 4  0.17248534          0.3021018 0.2098526 0.04403812
## 5  0.16164737          0.3072104 0.2160264 0.04666741
## 6  0.17266084          0.2816206 0.2121988 0.04502834
## 7  0.17390066          0.3081274 0.2112391 0.04462197
## 8  0.16222486          0.3506556 0.2134083 0.04554312
## 9  0.14615829          0.3145649 0.2160072 0.04665909
## 10 0.13639058          0.3258002 0.2195888 0.04821924
## 11 0.12781291          0.3624504 0.2171430 0.04715110
## 12 0.08839711          0.3752602 0.9515084 0.90536818
##
## [12 rows x 7 columns]
```

We are going to take as a initial point to create models the best 3 algorithm. In this case we will be comparing the Gradieng bosting machine, Generalized linear models and the eXtremely Randomized Trees.

```
gbm_model = h2o.gbm(x = predictor,y = "stroke",training_frame = train_h2o,keep_cross_validation_predict
```

## Gradieng Bosting Machine

```
##      |
```

The performance of this model is the follow.

```
gbm_perform = h2o.performance(gbm_model,newdata = test_h2o)
print(gbm_perform)
```

```
## H2OBinomialMetrics: gbm
##
## MSE:  0.03861112
## RMSE:  0.1964971
## LogLoss:  0.1405456
## Mean Per-Class Error:  0.3149004
## AUC:  0.8777287
## AUCPR:  0.2897502
## Gini:  0.7554574
## R^2:  0.1151304
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##          0  1  Error      Rate
## 0      898 62 0.064583  =62/960
## 1       26 20 0.565217  =26/46
## Totals 924 82 0.087475  =88/1006
##
## Maximum Metrics: Maximum metrics at their respective thresholds
##          metric threshold      value idx
## 1          max f1  0.145380  0.312500  73
```

```
## 2          max f2 0.055600 0.459057 178
## 3          max f0point5 0.595774 0.405405 6
## 4          max accuracy 0.595774 0.959245 6
## 5          max precision 0.746927 1.000000 0
## 6          max recall 0.015685 1.000000 312
## 7          max specificity 0.746927 1.000000 0
## 8          max absolute_mcc 0.595774 0.325158 6
## 9  max min_per_class_accuracy 0.055600 0.804348 178
## 10 max mean_per_class_accuracy 0.027221 0.820380 260
## 11          max tns 0.746927 960.000000 0
## 12          max fns 0.746927 45.000000 0
## 13          max fps 0.002404 960.000000 399
## 14          max tps 0.015685 46.000000 312
## 15          max tnr 0.746927 1.000000 0
## 16          max fnr 0.746927 0.978261 0
## 17          max fpr 0.002404 1.000000 399
## 18          max tpr 0.015685 1.000000 312
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/
```

```
glm_model = h2o.glm(x = predictor, y = "stroke", training_frame = train_h2o, keep_cross_validation_predictions = TRUE)
```

## Generalized Linear Models

```
## |
```

The performance of this model is the follow.

```
glm_perform = h2o.performance(glm_model, newdata = test_h2o)
print(glm_perform)
```

```
## H2OBinomialMetrics: glm
##
## MSE: 0.03901051
## RMSE: 0.1975108
## LogLoss: 0.1419672
## Mean Per-Class Error: 0.2157156
## AUC: 0.8741848
## AUCPR: 0.2108722
## Gini: 0.7483696
## R^2: 0.1059773
## Residual Deviance: 285.6381
## AIC: 329.6381
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##      0  1  Error  Rate
## 0    838 122 0.127083 =122/960
## 1     14  32 0.304348 =14/46
## Totals 852 154 0.135189 =136/1006
##
```

```

## Maximum Metrics: Maximum metrics at their respective thresholds
##
##      metric threshold      value idx
## 1      max f1  0.110358  0.320000 126
## 2      max f2  0.092543  0.492021 153
## 3      max f0point5 0.308445  0.265957  11
## 4      max accuracy 0.349515  0.955268   2
## 5      max precision 0.349515  0.666667   2
## 6      max recall  0.020681  1.000000 301
## 7      max specificity 0.356583  0.998958   0
## 8      max absolute_mcc 0.092543  0.341734 153
## 9      max min_per_class_accuracy 0.092543  0.804348 153
## 10     max mean_per_class_accuracy 0.092543  0.821445 153
## 11      max tns  0.356583 959.000000   0
## 12      max fns  0.356583  46.000000   0
## 13      max fps  0.001465 960.000000 399
## 14      max tps  0.020681  46.000000 301
## 15      max tnr  0.356583  0.998958   0
## 16      max fnr  0.356583  1.000000   0
## 17      max fpr  0.001465  1.000000 399
## 18      max tpr  0.020681  1.000000 301
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/

```

**Extremely Randomized Trees** This model is a type of random Forest who takes as many trees as predictors you have.

```
xrt_model = h2o.randomForest(x = predictor, y = "stroke", training_frame = train_h2o, keep_cross_validation
```

```
##      |
```

The performance of this model is the follow.

```
xrt_perform = h2o.performance(xrt_model, newdata = test_h2o)
print(xrt_perform)
```

```

## H2OBinoialMetrics: drf
##
## MSE:  0.04068915
## RMSE: 0.2017155
## LogLoss: 0.2363819
## Mean Per-Class Error: 0.2720788
## AUC:  0.8395833
## AUCPR: 0.2731995
## Gini: 0.6791667
## R^2: 0.06750706
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##      0  1  Error  Rate
## 0      855 105 0.109375  =105/960
## 1       20  26 0.434783  =20/46
## Totals 875 131 0.124254  =125/1006
##

```

```

## Maximum Metrics: Maximum metrics at their respective thresholds
##
##      metric threshold      value idx
## 1      max f1  0.160000  0.293785 40
## 2      max f2  0.120000  0.423729 43
## 3      max f0point5 0.480000  0.368852 14
## 4      max accuracy 0.788750  0.957256  4
## 5      max precision 0.936442  1.000000  0
## 6      max recall  0.000000  1.000000 61
## 7      max specificity 0.936442  1.000000  0
## 8      max absolute_mcc 0.480000  0.284252 14
## 9      max min_per_class_accuracy 0.060000  0.733333 53
## 10     max mean_per_class_accuracy 0.060000  0.779710 53
## 11      max tns  0.936442 960.000000  0
## 12      max fns  0.936442  45.000000  0
## 13      max fps  0.000000 960.000000 61
## 14      max tps  0.000000  46.000000 61
## 15      max tnr  0.936442  1.000000  0
## 16      max fnr  0.936442  0.978261  0
## 17      max fpr  0.000000  1.000000 61
## 18      max tpr  0.000000  1.000000 61
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/

```